

16 - System Design

CEG 4330/6330 - Microprocessor-Based Embedded Systems
Max Gilson

System Design

- In the lab you've been given a hardware platform and a set of requirements
 - Easy: we know the hardware platform can satisfy the requirements
- In embedded systems engineering you are only given requirements
 - Hard: what hardware satisfies these requirements?
 - Some systems may be impossible to build
- An embedded system engineer must be capable of designing either the entire system or parts of the system for integration

Challenges in Embedded System Design

- How much hardware or type of hardware?
 - How fast of a CPU? How much memory? I/O?
 - How much is too much?
 - E.g. \$500 CPU for a thermostat necessary?
- How can we meet deadlines?
 - Faster hardware or more clever software?
 - Is a RTOS required?
 - Might not be strict, customer may ask for a thermostat that sends a temperature reading once per minute
- How do we minimize power consumption?
 - Turn off unused components? Choose different hardware?

Before Starting the Design...

- Before starting our design we have to look at:
 - Requirements
 - The strict requirement that the design must satisfy
 - Some requirements may be impossible to meet!
 - E.g. a smartphone that never dies and doesn't need to be plugged in

After Completing the Design...

- After completing our design we have to look at:
 - Specification
 - What our design “says” it’s capable of
 - Implementation
 - Our actual physical implementation, hardware, software, and all components

After Completing the Design... (cont.)

- Does our design actually work?
 - Are the requirements realistic and what the customer really wanted?
 - Does the specification actually satisfy the requirements?
 - Does the implementation meet the specification?
 - How do we test the device in real time?
 - When you have a prototype or the final product how can we make sure it does what it says?
 - How do we test using real data?
 - E.g. how do we ensure the thermostat reads temperature and sets temperatures properly?

After Completing the Design... (cont.)

- How do we work with our system?
 - Observability
 - How can we observe what's going on inside the system?
 - Battery voltage reading? Data inside memory?
 - Controllability
 - How can we control our device?
 - What is our development platform?
 - Ex: Arduino IDE, Raspberry Pi OS, etc.

Design Methodologies


- A design methodology is a process for designing the system
 - Sometimes called a workflow
- It is important to have a process so that no steps are skipped
- Some software tools can help automate steps in the process or keep track of the process itself
 - Compilers, Jira/Trello, CAD tools, etc.
 - Your boss, customers, or certification organizations might need to see you've tracked your process accurately

Design Goals

- Your requirements are most likely going to involve one or more of these goals:
 - Performance
 - Functionality and user interface
 - Manufacturing cost
 - Power consumptions
 - Other requirements
 - Physical size, weight, durability, water submerging, etc.


Layers of Design Methodology

- Requirement
 - An outline of what the end product needs to be capable of
 - Example:
 - A thermostat that reads and controls temperature
 - Connects to your phone
 - Can be controlled via the Internet




Layers of Design Methodology

- Specification
 - A technical outline that will meet the requirements
 - The specification is for the engineers to follow and adhere to
 - Example:
 - The thermostat will be powered off of 120V AC
 - The temperature reading should have a range of 50F to 90F and an accuracy of ± 1 F
 - Wireless internet connectivity



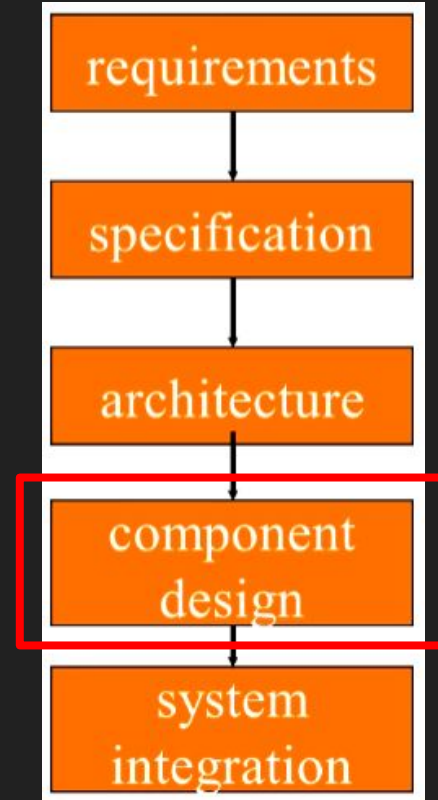
Layers of Design Methodology

- Architecture
 - The high level connectivity of all the major pieces
 - Example:
 - A voltage regulator section to convert 120V AC to DC
 - A microcontroller to read and control temperature
 - A PID loop in software needed for temperature control
 - A WiFi section




Layers of Design Methodology

- Component Design
 - The selection of the actual components
 - Example:
 - Selecting all of the resistors, capacitors, regulators for converting 120V AC to DC
 - An Atmega328p for the microcontroller
 - A u-blox NINA-W102 for WiFi receiver




Layers of Design Methodology

- System Integration
 - The actual physical combination of all the components
 - Example:
 - The manufactured circuit board housing all the components and with software installed



Top Down vs Bottom Up

- Top Down Design
 - Start with requirements
 - Build out each step along the way
 - Ensures no steps are skipped and requirements are met
- Bottom Up Design
 - Start with a component(s) then build out into bigger system
 - Could involve development boards or reverse engineering
 - Demonstrates hard to measure values like power consumption, memory usage, speed, etc.
- In practice, both approaches are used



Where to Get Requirements

- Requirements are plain language descriptions of what the user wants and expects
- Requirements can be developed from:
 - Talking directly to your customers
 - Talking to marketing representatives or doing your own market research
 - Providing prototypes to users for comment

Functional vs Non-Functional Requirements

- Functional requirements:
 - Output as a function of input
- Non-functional requirements:
 - Performance
 - Size and/or weight
 - Power consumption
 - Reliability
 - Etc.

Design Process Example (Requirements)

- The marketing team at your company asks you (the engineer) to build a new product
- The product is a dash cam requiring:
 - A “real-time” camera
 - Powered by a car battery
 - Holds onto the last 2 hours of video
 - Mounts to a windshield



Design Process Example (Specifications)

- We can establish some specifications from the requirements
- Requirement: A “real-time” camera
 - Spec: A 480p 25 FPS camera controlled via microcontroller
 - Basis: Seems attainable based on some bottom up design (more later)
- Requirement: Powered by a car battery
 - Spec: Input voltage of 15V to 12V DC
 - Basis: A car’s battery voltage ranges from 15V to 12V DC
- Requirement: Holds onto the last 2 hours of video
 - Spec: 4 GB SD card
 - Basis: <https://www.dr-lex.be/info-stuff/videocalc.html>
- Requirement: Mounts to a windshield
 - Spec: A special case design that encloses the electronics and has a suction cup
 - Basis: A suction cup is a great mount for a windshield

Design Process Example (Architecture)

- Next, we can create a block diagram to outline our architecture
- We will most likely need the following blocks based on our specification:
 - Camera
 - Microcontroller or Microprocessor
 - Source code outline
 - Power regulation
 - SD card

Design Process Example (Component Design)

- Some bottom up design may prove useful in deciding our components
- First, we can survey different dash cams (reverse engineering) and camera systems to see if there already exists an open source or “solved” system
 - The ESP32-CAM satisfies the “real-time” camera requirement and can run at 480p at 25 FPS
 - Source:
<https://github.com/jameszah/ESP32-CAM-Video-Recorder-junior>
- Next, we can start breaking down our architecture into components that make them up
 - Microcontroller circuits
 - Power regulation circuits
 - Source code

Design Process Example (Implementation)

- Next is building the actual implementation
 - Purchase the components and assemble it yourself
OR ask a manufacturer to assemble it for you
 - Write and install any software required (easier said than done)
 - Test the implementation and verify that it meets the specification
 - If yes: success!
 - If no: iterate, fix issues, revise the architecture or components, etc.

Design Verification

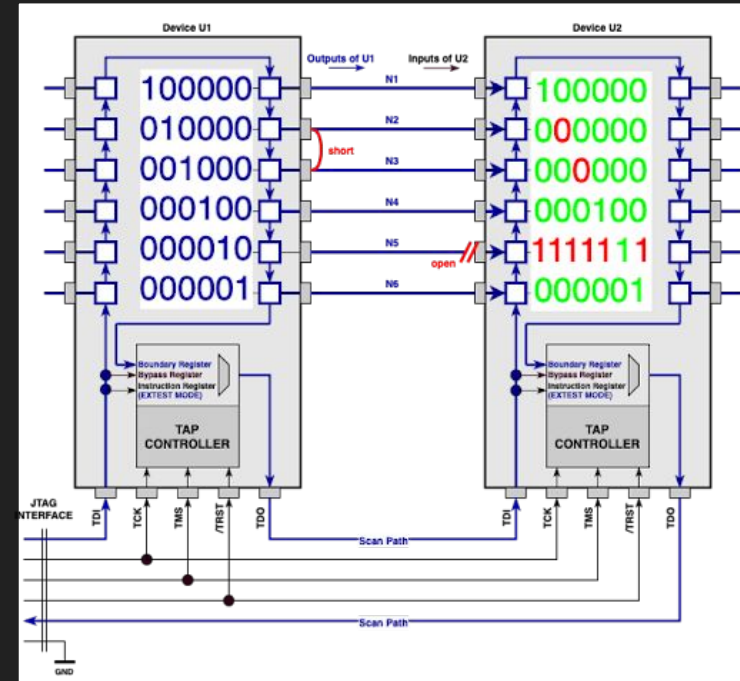

- It is important to make sure the design meets the specification
 - You should verify your design
- If you have not tested it how do you know?
 - For example, if we say:
“Our device can take an input voltage from 15V to 12V DC.”
Does the device always work in this range? Is there performance degradation at 12V vs 15V?

Built-In Self Test (BIST)

- A built-in self test is a great method for verifying the design works as intended
 - For example, if you built a drone, you may want it to check some things when it is powered up
 - Verify the microcontroller can communicate with sensors
 - Verify battery is healthy
 - Verify memory with a checksum
- This can be done by a built-in test that automatically runs
 - This concept applies not just to embedded systems, but any system
 - Example: could apply to the anti-lock brake system in a car

JTAG Testing

- Joint Test Action Group (JTAG) is an industry standard method for testing hardware
 - Replaced bed of nails testing
- JTAG is hardware inside the microcontroller itself that allows you to programm, debug, and verify pin functionality inside of the microcontroller



The Inevitable Moving Goal Posts

- After you complete your design or prototype, the marketing team might say:
 - “We can’t sell a 480p dash cam! The quality is too bad! Let’s revise the requirements to require 1080p!”
- Do we have to change anything?
 - Hint: yes, almost everything
- Moral of the story: make sure you have ALL the requirements before you design

Designing for Reliability

- How can we quantify if a system is reliable or not?
 - Mean Time to Failure (MTTF)
 - Mean Time Between Failures (MTBF)
 - Availability is equal to $MTTF / MTBF$
- There exists some probability that a system will fail
 - This probability changes over time
- Repairing failures can be done by:
 - Mask errors (hide the error)
 - Hot spare (swap with spare part)

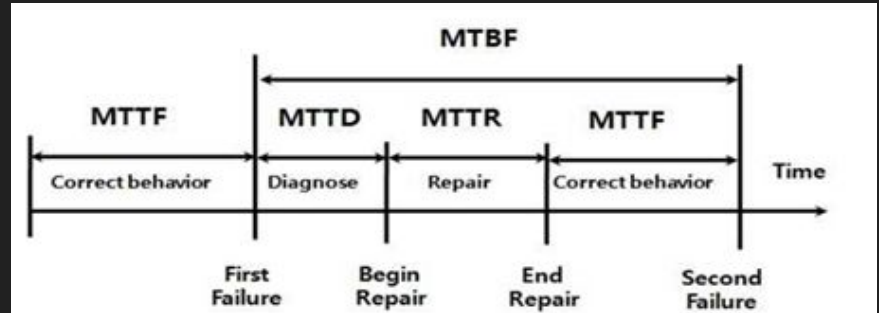


Fig. 1 MTDD, MTTR, MTTF & MTBF





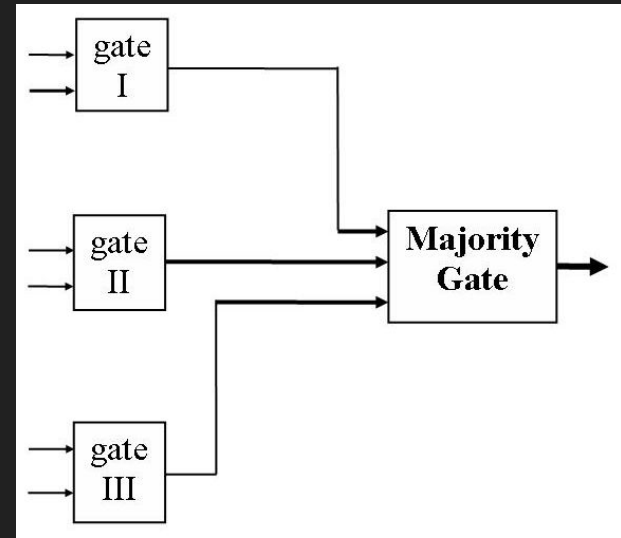
Fig. 2: Bathtub curve illustrate consistent rate of failures

DMR and TMR

- Dual modular redundancy (DMR) is having two copies of the same hardware or software component working in parallel
 - Software components should be developed by separate teams using the same specification
 - Provides error detection but not error correction
 - If my components both give me a different output, which is correct?

DMR and TMR (cont.)

- Triple modular redundancy (TMR) is having three copies of the same hardware or software component working in parallel
 - Provides error detection AND error correction
 - When using 3 parallel components we can perform error correction
 - If 1 component is outputting something different than the other 2 it can be replaced
 - A gate can be developed to take the majority of outputs
- $R_{TMR} = R_v (3 R_m^2 - 2 R_m^3)$
 - R_v is probability the majority gate (voter) fails
 - R_m is probability the node fails
 - R_{TMR} is the probability the whole system fails



ECC Memory

- Error Checking and Correction (ECC) is a method for verifying that memory is operating properly
 - A memory chip might have 64 data pins or 72 data pins
 - If using 72 pins: 64 are data bits, 8 are check bits
- Using a Hamming code encoder, there are 4 data bits and 3 check bits, then we have 3 even-parity equations
 - Check bits are produced to an even parity for each equation
 - If any of the equations output non-zero, an error is detected and can be corrected
 - If output is (1, 1, 1) then D3 must be wrong
 - If output is (1, 0, 0) then C2 must be wrong
 - Fails when two bits are wrong (but this is unlikely)

C2	\oplus	D3	\oplus	D2	\oplus	D1		= 0
C1	\oplus	D3	\oplus	D2			\oplus D0	= 0
C0	\oplus	D3			\oplus	D1	\oplus D0	= 0

Burst Errors

- Hamming code is useful when a noisy environment might corrupt communication
 - A major downfall is we can only detect 1 bit errors
- Burst errors are going to cause a burst of errors to occur in short periods of time
 - Can result in multiple bits being wrong
- Reed-Solomon error correction can mitigate this issue but is significantly more complicated

Cyclic Redundancy Check (CRC)

- A cyclic redundancy check module can create a signature for any piece of data sent
- This signature can be recreated on both the receiver and the transmitter side of the communication
- If there is an inconsistent signature on either side, the data received contained errors
- Any size of data can be sent either 1 bit or 1000's of bits
 - A signature may not always be unique

